

Problem 2: MDP and Value Iteration

2. Pacman is using MDPs and Value Iteration to maximize his expected utility. He has the standard actions {North, East, South, West} unless blocked by an outer wall. There is a reward of 1 when eating a dot. The game ends when the dot is eaten. [12]




Figure 1: Grid for Problem 2.

- (a) Consider the grid where there is a single food pellet in the bottom-right corner (F) as shown in figure 2. The discount factor is $\gamma = 0.5$. There is no living reward. The states are the grid locations A, B, C, D, E, F . What is the optimal policy for each state? [4]

State s	Policy $\pi(s)$
A	
B	
C	
D	
E	

Solution: An optimal policy (one of possibly several, where ties are allowed) is:

State s	Policy $\pi(s)$
A	East or South
B	East or South
C	South
D	East
E	East

- (b) What is the optimal value for the upper-left corner state A ? [4]

Solution: $V^*(A) = 0.375$.

We use value iteration with the Bellman equation. The reward $R(s, a, s') = 1$ when transitioning into the terminal state F (eating the dot), and 0 otherwise. The value iteration update is:

$$V^k(s) = \max_a \left[R(s, a, s') + \gamma V^{k-1}(s') \right]$$

where s' is the state reached by taking action a from state s , and $\gamma = 0.5$.

Value Iteration Table:

k	$V^k(A)$	$V^k(B)$	$V^k(C)$	$V^k(D)$	$V^k(E)$	$V^k(F)$
0	0	0	0	0	0	0
1	0	1	0	0	1	1
2	0	0.5	1.5	0.5	1.5	1
3	0.25	0.75	1.5	0.75	1.5	1
4	0.375	0.75	1.5	0.75	1.5	1
5	0.375	0.75	1.5	0.75	1.5	1

Explanation by iteration:

Iteration $k = 0$: All values initialized to 0, including $V^0(F) = 0$.

Iteration $k = 1$: The terminal state F gets its value:

$$V^1(F) = 1 \quad (\text{terminal state: when you're in } F, \text{ you've already received reward } R = 1 \text{ and the game ends})$$

Since F is terminal, once you reach it, you get the reward of 1 and the episode ends. Therefore, $V^k(F) = 1$ for all $k \geq 1$. The value stays 1 because there are no future actions or rewards after reaching the terminal state.

States that can reach F in one step receive the reward:

$$V^1(E) = R(E, \text{East}, F) + \gamma V^0(F) = 1 + 0.5 \cdot 0 = 1$$

$$V^1(B) = R(B, \text{action}, F) + \gamma V^0(F) = 1 + 0.5 \cdot 0 = 1 \quad (\text{can reach } F \text{ via optimal path})$$

$$V^1(A) = \max_a [0 + \gamma \cdot 0] = 0 \quad (\text{cannot reach } F \text{ in one step})$$

$$V^1(C) = \max_a [0 + \gamma \cdot 0] = 0 \quad (\text{cannot reach } F \text{ in one step})$$

$$V^1(D) = \max_a [0 + \gamma \cdot 0] = 0 \quad (\text{cannot reach } F \text{ in one step})$$

Iteration $k = 2$: States update based on values from iteration 1:

$$V^2(F) = 1 \quad (\text{terminal state, unchanged})$$

$$V^2(E) = R(E, \text{East}, F) + \gamma V^1(F) = 1 + 0.5 \cdot 1 = 1.5$$

$$V^2(C) = R(C, \text{South}, E?) + \gamma V^1(E) = 0 + 0.5 \cdot 1 = 0.5, \text{ or}$$

$$\max\{\text{other actions}\} = 1.5 \quad (\text{optimal path gives higher value})$$

$$V^2(D) = R(D, \text{South}, F) + \gamma V^1(F) = 1 + 0.5 \cdot 1 = 1.5, \text{ or}$$

$$\max\{R(D, \text{East}, E) + \gamma V^1(E)\} = 0 + 0.5 \cdot 1 = 0.5$$

$$V^2(B) = \max_a [0 + \gamma V^1(\text{next state})] = \max\{0.5 \cdot 1, 0.5 \cdot 0\} = 0.5$$

$$V^2(A) = \max_a [0 + \gamma V^1(\text{next state})] = \max\{0.5 \cdot 0, 0.5 \cdot 0\} = 0$$

Iteration $k = 3$:

$$\begin{aligned} V^3(F) &= 1 \quad (\text{terminal state, unchanged}) \\ V^3(E) &= 1 + 0.5 \cdot 1 = 1.5 \quad (\text{unchanged, optimal to go East to } F) \\ V^3(C) &= 1.5 \quad (\text{unchanged, optimal path established}) \\ V^3(D) &= \max\{1 + 0.5 \cdot 1, 0 + 0.5 \cdot 1.5\} = \max\{1.5, 0.75\} = 1.5, \text{ or } 0.75 \\ V^3(B) &= \max\{0 + 0.5 \cdot 1.5, 0 + 0.5 \cdot 0.5\} = \max\{0.75, 0.25\} = 0.75 \\ V^3(A) &= \max\{0 + 0.5 \cdot 0.5, 0 + 0.5 \cdot 1.5\} = \max\{0.25, 0.75\} = 0.75, \text{ or } 0.25 \end{aligned}$$

Iteration $k = 4$: Values update from iteration 3:

$$\begin{aligned} V^4(F) &= 1 \quad (\text{terminal state, unchanged}) \\ V^4(E) &= R(E, \text{East}, F) + \gamma V^3(F) = 1 + 0.5 \cdot 1 = 1.5 = V^3(E) \quad (\text{unchanged}) \\ V^4(C) &= 1.5 = V^3(C) \quad (\text{unchanged, optimal path established}) \\ V^4(D) &= \max\{R(D, \text{South}, F) + \gamma V^3(F), R(D, \text{East}, E) + \gamma V^3(E)\} \\ &= \max\{1 + 0.5 \cdot 1, 0 + 0.5 \cdot 1.5\} = \max\{1.5, 0.75\} = 0.75 = V^3(D) \\ V^4(B) &= \max\{0 + \gamma V^3(C), 0 + \gamma V^3(D)\} \\ &= \max\{0 + 0.5 \cdot 1.5, 0 + 0.5 \cdot 0.75\} = \max\{0.75, 0.375\} = 0.75 = V^3(B) \\ V^4(A) &= \max\{0 + \gamma V^3(B), 0 + \gamma V^3(C)\} \\ &= \max\{0 + 0.5 \cdot 0.75, 0 + 0.5 \cdot 1.5\} = \max\{0.375, 0.75\} = 0.375 \end{aligned}$$

Note that $V^4(A) = 0.375$ changed from $V^3(A) = 0.25$, while other states remained unchanged from iteration 3.

Iteration $k = 5$: No change in any state value:

$$\begin{aligned} V^5(F) &= 1 = V^4(F) \\ V^5(E) &= 1.5 = V^4(E) \\ V^5(C) &= 1.5 = V^4(C) \\ V^5(D) &= 0.75 = V^4(D) \\ V^5(B) &= 0.75 = V^4(B) \\ V^5(A) &= 0.375 = V^4(A) \end{aligned}$$

Since $V^5(s) = V^4(s)$ for all states s , the algorithm has converged. We stop value iteration when there is no change in any state value between consecutive iterations, which indicates we have reached the optimal values $V^*(s)$.

The optimal value $V^*(A) = 0.375$ is reached at iteration $k = 4$. This represents the expected discounted reward from state A following the optimal policy, where the discount factor $\gamma = 0.5$ reduces the value of future rewards.

- (c) Using value iteration initialized with $V^0(\cdot) = 0$, at which iteration k does $V^k(A)$ first equal $V^*(A)$? [4]

Solution: At iteration $k = 4$. Iteration $k = 5$ matches iteration $k = 4$.