

## Homework 2

### Instructions


*Points:* Please see the points for each problem.

*Submission Instructions:* Please submit a PDF in Canvas.

### Points

Question	Part	Points Possible	Points Earned
1. Minimax and $\alpha$ - $\beta$ pruning	(a)	4	_____
	(b)	4	_____
2. MDP and Value Iteration	(a)	4	_____
	(b)	4	_____
	(c)	4	_____
3. Policy Gradient and PPO	(a)	3	_____
	(b)	3	_____
<b>Total</b>		<b>26</b>	_____

1. Minimax and  $\alpha$ - $\beta$  pruning Consider the minimax tree shown below for parts (a) and (b).



**Figure 1:** Minimax tree for Question 1.

- (a) What value will the root node *A* have?

**Solution:** From the leaves, compute internal node values bottom-up:


$$\begin{aligned} D &= \max\{15, 6, 21\} = 21, \\ B &= \min\{12, D\} = 12, \\ E &= \max\{5, 10\} = 10, \\ F &= \max\{1, 19\} = 19, \\ C &= \min\{E, F\} = 10, \\ A &= \max\{B, C, 2\} = 12. \end{aligned}$$

Thus, the root value is  $A = 12$ .

- (b) Cross off the nodes that are pruned by  $\alpha$ - $\beta$  pruning. Assume standard left-to-right traversal. If a non-terminal state ( $A, B, C, D, E$ , or  $F$ ) is pruned, cross off the entire subtree.

**Solution:** With left-to-right traversal, the leaf with value 6 and the leaf with value 21 under  $D$  are pruned, and the entire subtree under  $F$  is pruned.

2. Pacman is using MDPs and Value Iteration to maximize his expected utility. He has the standard actions {North, East, South, West} unless blocked by an outer wall. There is a reward of 1 when eating a dot. The game ends when the dot is eaten.



**Figure 2:** Grid for Question 2.

- (a) Consider the grid where there is a single food pellet in the bottom-right corner ( $F$ ) as shown in figure 2. The discount factor is  $\gamma = 0.5$ . There is no living reward. The states are the grid locations  $A, B, C, D, E, F$ . What is the optimal policy for each state?

State $s$	Policy $\pi(s)$
$A$	
$B$	
$C$	
$D$	
$E$	

**Solution:** An optimal policy (one of possibly several, where ties are allowed) is:

State $s$	Policy $\pi(s)$
$A$	East or South
$B$	East or South
$C$	South
$D$	East
$E$	East

- (b) What is the optimal value for the upper-left corner state  $A$ ?

**Solution:**  $V^*(A) = 0.375$ .

We use value iteration with the Bellman equation. The reward  $R(s, a, s') = 1$  when transitioning into the terminal state  $F$  (eating the dot), and 0 otherwise. The value iteration update is:

$$V^k(s) = \max_a [R(s, a, s') + \gamma V^{k-1}(s')]$$

where  $s'$  is the state reached by taking action  $a$  from state  $s$ , and  $\gamma = 0.5$ .

**Value Iteration Table:**

$k$	$V^k(A)$	$V^k(B)$	$V^k(C)$	$V^k(D)$	$V^k(E)$	$V^k(F)$
0	0	0	0	0	0	0
1	0	1	0	0	1	1
2	0	0.5	1.5	0.5	1.5	1
3	0.25	0.75	1.5	0.75	1.5	1
4	0.375	0.75	1.5	0.75	1.5	1
5	0.375	0.75	1.5	0.75	1.5	1

**Explanation by iteration:**

**Iteration  $k = 0$ :** All values initialized to 0, including  $V^0(F) = 0$ .

**Iteration  $k = 1$ :** The terminal state  $F$  gets its value:

$$V^1(F) = 1 \quad (\text{terminal state: when you're in } F, \text{ you've already received reward } R = 1 \text{ and the game ends})$$

Since  $F$  is terminal, once you reach it, you get the reward of 1 and the episode ends. Therefore,  $V^k(F) = 1$  for all  $k \geq 1$ . The value stays 1 because there are no future actions or rewards after reaching the terminal state.

States that can reach  $F$  in one step receive the reward:

$$V^1(E) = R(E, \text{East}, F) + \gamma V^0(F) = 1 + 0.5 \cdot 0 = 1$$

$$V^1(B) = R(B, \text{action}, F) + \gamma V^0(F) = 1 + 0.5 \cdot 0 = 1 \quad (\text{can reach } F \text{ via optimal path})$$

$$V^1(A) = \max_a [0 + \gamma \cdot 0] = 0 \quad (\text{cannot reach } F \text{ in one step})$$

$$V^1(C) = \max_a [0 + \gamma \cdot 0] = 0 \quad (\text{cannot reach } F \text{ in one step})$$

$$V^1(D) = \max_a [0 + \gamma \cdot 0] = 0 \quad (\text{cannot reach } F \text{ in one step})$$

**Iteration  $k = 2$ :** States update based on values from iteration 1:

$$\begin{aligned}
 V^2(F) &= 1 \quad (\text{terminal state, unchanged}) \\
 V^2(E) &= R(E, \text{East}, F) + \gamma V^1(F) = 1 + 0.5 \cdot 1 = 1.5 \\
 V^2(C) &= R(C, \text{South}, E?) + \gamma V^1(E) = 0 + 0.5 \cdot 1 = 0.5, \text{ or} \\
 &\quad \max\{\text{other actions}\} = 1.5 \quad (\text{optimal path gives higher value}) \\
 V^2(D) &= R(D, \text{South}, F) + \gamma V^1(F) = 1 + 0.5 \cdot 1 = 1.5, \text{ or} \\
 &\quad \max\{R(D, \text{East}, E) + \gamma V^1(E)\} = 0 + 0.5 \cdot 1 = 0.5 \\
 V^2(B) &= \max_a [0 + \gamma V^1(\text{next state})] = \max\{0.5 \cdot 1, 0.5 \cdot 0\} = 0.5 \\
 V^2(A) &= \max_a [0 + \gamma V^1(\text{next state})] = \max\{0.5 \cdot 0, 0.5 \cdot 0\} = 0
 \end{aligned}$$

**Iteration  $k = 3$ :**

$$\begin{aligned}
 V^3(F) &= 1 \quad (\text{terminal state, unchanged}) \\
 V^3(E) &= 1 + 0.5 \cdot 1 = 1.5 \quad (\text{unchanged, optimal to go East to } F) \\
 V^3(C) &= 1.5 \quad (\text{unchanged, optimal path established}) \\
 V^3(D) &= \max\{1 + 0.5 \cdot 1, 0 + 0.5 \cdot 1.5\} = \max\{1.5, 0.75\} = 1.5, \text{ or } 0.75 \\
 V^3(B) &= \max\{0 + 0.5 \cdot 1.5, 0 + 0.5 \cdot 0.5\} = \max\{0.75, 0.25\} = 0.75 \\
 V^3(A) &= \max\{0 + 0.5 \cdot 0.5, 0 + 0.5 \cdot 1.5\} = \max\{0.25, 0.75\} = 0.75, \text{ or } 0.25
 \end{aligned}$$

**Iteration  $k = 4$ :** Values update from iteration 3:

$$\begin{aligned}
 V^4(F) &= 1 \quad (\text{terminal state, unchanged}) \\
 V^4(E) &= R(E, \text{East}, F) + \gamma V^3(F) = 1 + 0.5 \cdot 1 = 1.5 = V^3(E) \quad (\text{unchanged}) \\
 V^4(C) &= 1.5 = V^3(C) \quad (\text{unchanged, optimal path established}) \\
 V^4(D) &= \max\{R(D, \text{South}, F) + \gamma V^3(F), R(D, \text{East}, E) + \gamma V^3(E)\} \\
 &= \max\{1 + 0.5 \cdot 1, 0 + 0.5 \cdot 1.5\} = \max\{1.5, 0.75\} = 0.75 = V^3(D) \\
 V^4(B) &= \max\{0 + \gamma V^3(C), 0 + \gamma V^3(D)\} \\
 &= \max\{0 + 0.5 \cdot 1.5, 0 + 0.5 \cdot 0.75\} = \max\{0.75, 0.375\} = 0.75 = V^3(B) \\
 V^4(A) &= \max\{0 + \gamma V^3(B), 0 + \gamma V^3(C)\} \\
 &= \max\{0 + 0.5 \cdot 0.75, 0 + 0.5 \cdot 1.5\} = \max\{0.375, 0.75\} = 0.375
 \end{aligned}$$

Note that  $V^4(A) = 0.375$  changed from  $V^3(A) = 0.25$ , while other states remained unchanged from iteration 3.

**Iteration  $k = 5$ :** No change in any state value:

$$\begin{aligned} V^5(F) &= 1 = V^4(F) \\ V^5(E) &= 1.5 = V^4(E) \\ V^5(C) &= 1.5 = V^4(C) \\ V^5(D) &= 0.75 = V^4(D) \\ V^5(B) &= 0.75 = V^4(B) \\ V^5(A) &= 0.375 = V^4(A) \end{aligned}$$

Since  $V^5(s) = V^4(s)$  for all states  $s$ , the algorithm has converged. We stop value iteration when there is no change in any state value between consecutive iterations, which indicates we have reached the optimal values  $V^*(s)$ .

The optimal value  $V^*(A) = 0.375$  is reached at iteration  $k = 4$ . This represents the expected discounted reward from state  $A$  following the optimal policy, where the discount factor  $\gamma = 0.5$  reduces the value of future rewards.

- (c) Using value iteration initialized with  $V^0(\cdot) = 0$ , at which iteration  $k$  does  $V^k(A)$  first equal  $V^*(A)$ ?

**Solution:** At iteration  $k = 4$ . Iteration  $k = 5$  matches iteration  $k = 4$ .

3. Consider some of the newer RL algorithms such as Policy Gradient and Proximal Policy Optimization.

- (a) What issues with RL are they designed to solve?

**Solution:** Policy Gradient and PPO algorithms address several key issues in RL:

- **Sample inefficiency:** Traditional value-based methods (like Q-learning) require many samples to learn good policies, especially in high-dimensional or continuous action spaces.
- **High variance:** Policy gradient methods suffer from high variance in gradient estimates, making training unstable.
- **Stability:** Vanilla policy gradients can make large, destructive policy updates that degrade performance.
- **On-policy data requirement:** Many RL algorithms require on-policy data, making sample reuse difficult and inefficient.
- **Continuous action spaces:** Value-based methods struggle with continuous actions, while policy gradients handle them naturally.

*Note: Stating any 3 of these is sufficient for full credit.*

- (f) How do they improve on previous RL algorithms?

**Solution:** Policy Gradient and PPO improve on previous RL algorithms in several ways:

- **Direct policy optimization:** Policy gradients optimize the policy directly, avoiding the need to learn value functions first (unlike value-based methods like Q-learning).
- **PPO clipping:** PPO uses a clipped objective function that prevents large policy updates, maintaining training stability while allowing multiple updates from the same data.
- **Better sample efficiency:** PPO can perform multiple gradient updates on the same batch of data, improving sample efficiency compared to vanilla policy gradients.
- **Natural continuous actions:** Policy gradients parameterize policies directly, making them well-suited for continuous action spaces without discretization.
- **Reduced variance:** PPO's clipped objective and trust region approach reduce variance compared to standard policy gradient methods.
- **Stable learning:** The clipping mechanism ensures policy updates stay within a trust region, preventing performance collapse that can occur with large updates.

*Note: Stating any 3 of these is sufficient for full credit.*