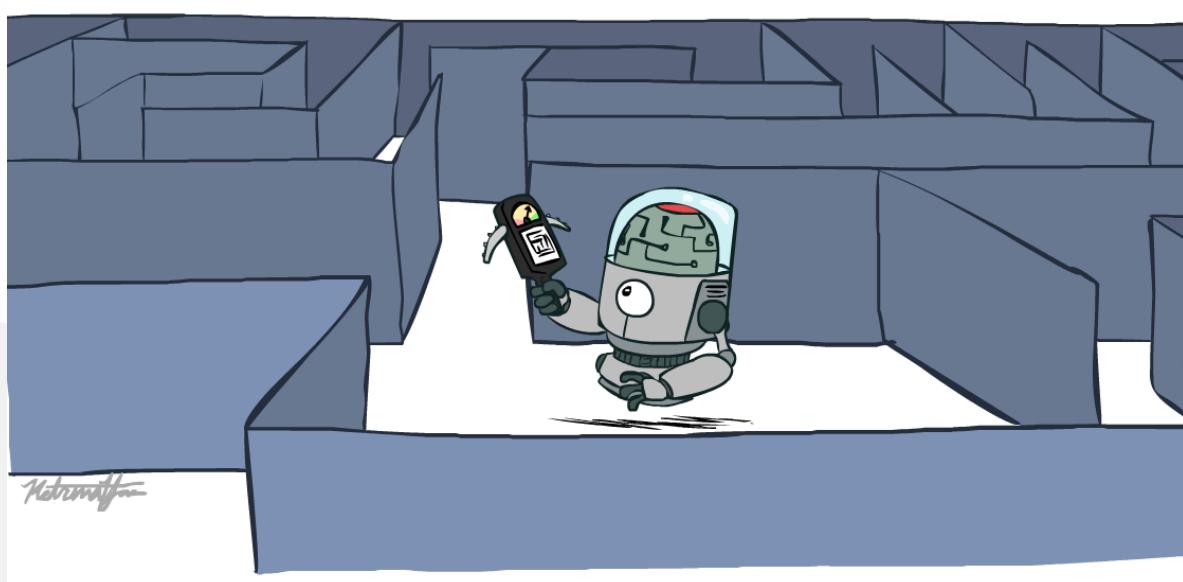


CS 6460: Artificial Intelligence

Informed Search



Instructor: George Rudolph
Utah Valley University Spring 2025

[These slides adapted from Dan Klein and Pieter Abbeel at UC Berkley]

Learning Outcomes

1. Solve Problems using Informed Searches

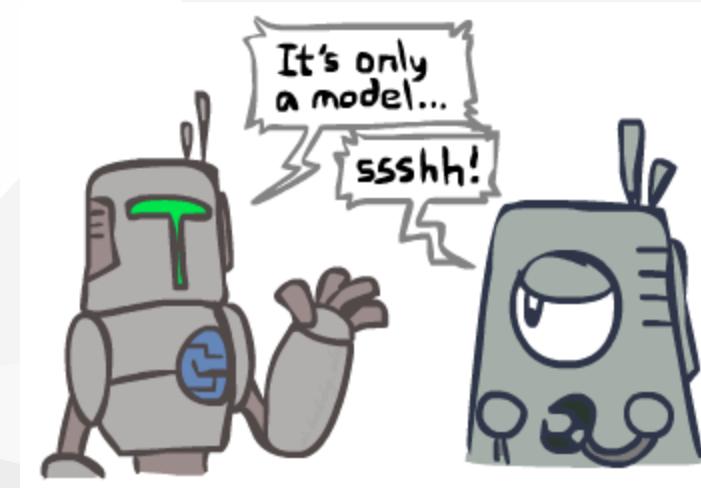
- Heuristics
- Greedy Search
- A* Search

2. Model Problems as Graph Search



Search and Models

- Search operates over models of the world
- The agent doesn't actually try all the plans out in the real world!
- Planning is all **in simulation**
- Your search is only as good as your models...



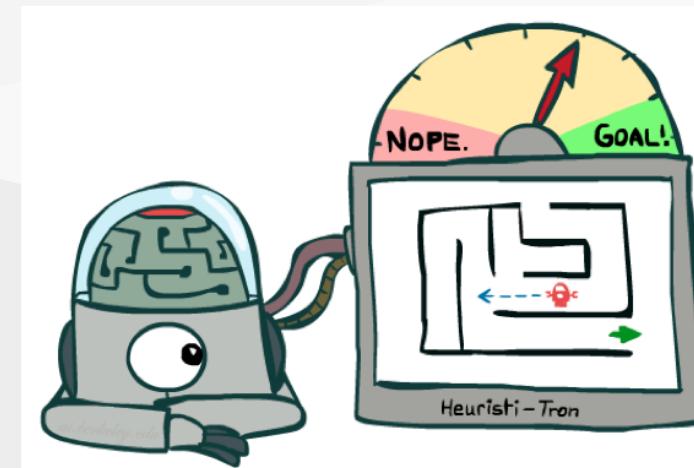
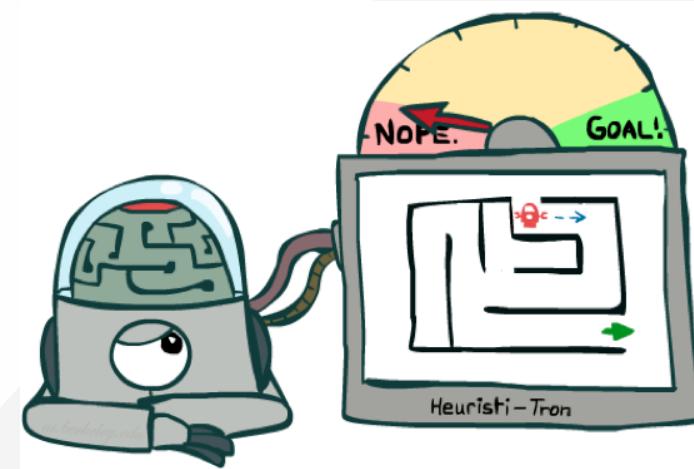
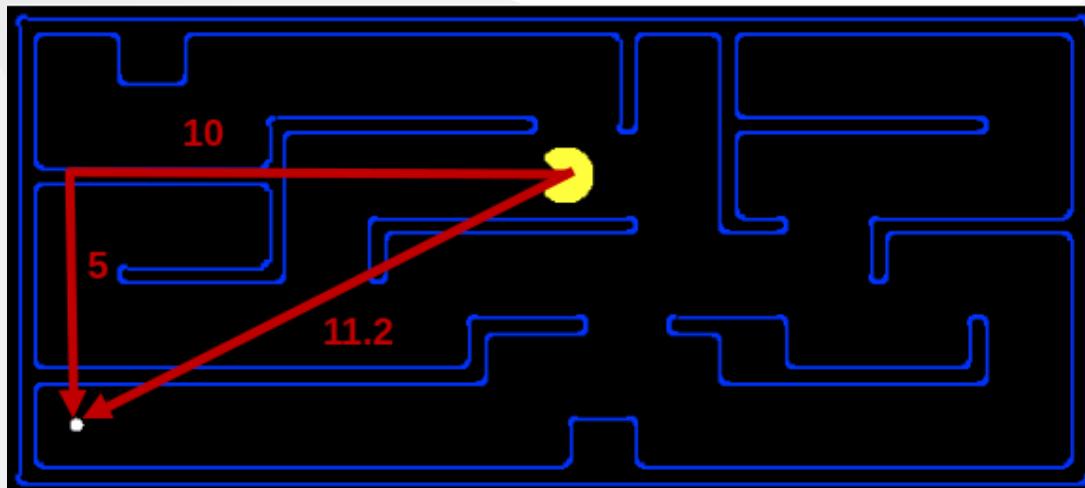
Informed Search



Search Heuristics

A heuristic is:

- A function that **estimates** how close a state is to a goal
- Designed for a **particular** search problem
- Examples: Manhattan distance, Euclidean distance



Example: Heuristic Function

- $h(x)$

 Example_Heuristic_Function

Converted shape

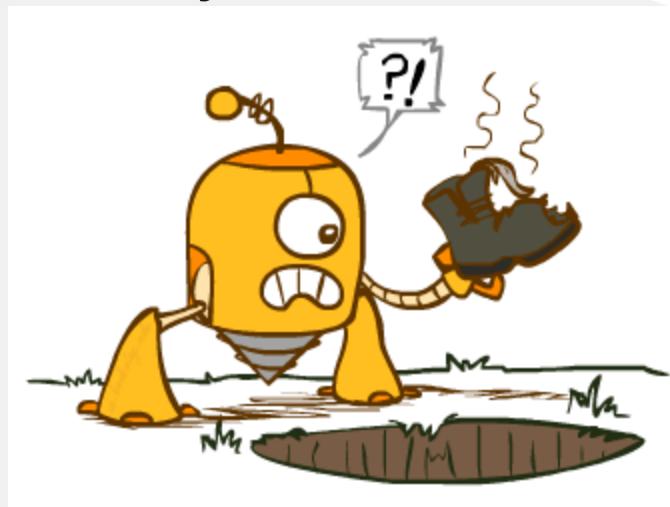
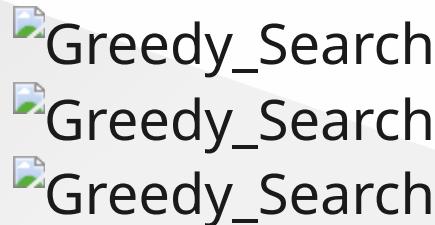
$h(x)$

Greedy Search



Greedy Search

- Expand the node that seems closest...
- What can go wrong?



Greedy Search

- Strategy: expand a node that you think is closest to a goal state
- Heuristic: estimate of distance to nearest goal for each state
- A common case:
- Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS
- ...
- b
- ...
- b

Converted shape

...

Video of Demo Contours Greedy (Empty)

Video of Demo Contours Greedy (Pacman Small Maze)

A* Search



Combining UCS and Greedy

- Uniform-cost orders by path cost, or backward cost $g(n)$
- Greedy orders by goal proximity, or forward cost $h(n)$
- A* Search orders by the sum: $f(n) = g(n) + h(n)$
- S
- a
- d
- b
- G
- $h=5$
- $h=6$
- $h=2$
- 1
- 8

When should A* terminate?

- Should we stop when we put a goal in the fringe?
- No: only stop when we pull a goal off the fringe
- S
- B
- A
- G
- 2
- 3
- 2
- 2
- $h = 1$
- $h = 2$
- $h = 0$

Is A* Optimal?

- What is wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!
- A
- G
- S
- 1
- 3
- $h = 6$
- $h = 0$
- 5
- $h = 7$
- Converted shape

Admissible Heuristics



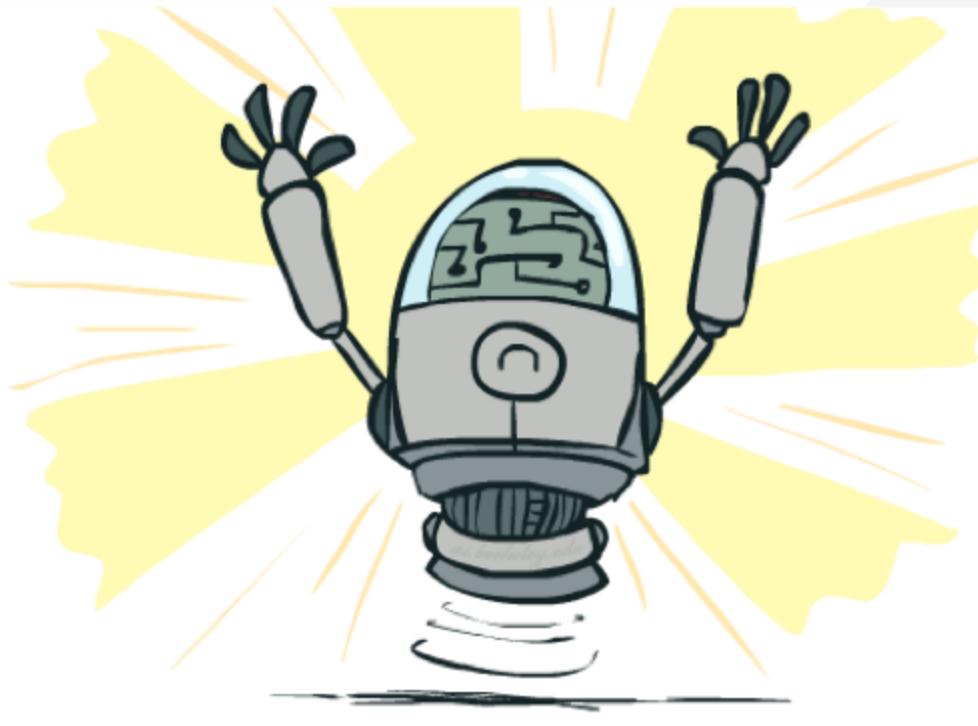
Admissible Heuristics

- A heuristic h is admissible (optimistic) if:
- where is the true cost to a nearest goal
- Examples:
- Defining admissible heuristics is the biggest effort in using A* in practice

$$0 \leq h(n) \leq h^*(n)$$

$h^*(n)$

Optimality of A* Tree Search



Optimality of A* Tree Search

- Assume:

Optimality of A* Tree Search: Blocking

- Proof:
- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
- $f(n)$ is less or equal to $f(A)$
- Definition of f-cost
- Admissibility of h
- ...
- $h = 0$ at a goal

$$f(n) = g(n) + h(n)$$

Optimality of A* Tree Search: Blocking

- Proof:
- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
- $f(n)$ is less or equal to $f(A)$
- $f(A)$ is less than $f(B)$
- B is suboptimal
- $h = 0$ at a goal
- ...

$$f(A) < f(B)$$

Optimality of A* Tree Search: Blocking

- Proof:
- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
- $f(n)$ is less or equal to $f(A)$
- $f(A)$ is less than $f(B)$
- n expands before B
- All ancestors of A expand before B
- A expands before B
- A* search is optimal
- ...

$$f(n) < f(A) < f(B)$$

Properties of A*

Properties of A*

- ...
- b
- ...
- b
- Uniform-Cost
- A*

Converted shape

...

Converted shape

b

UCS vs A* Contours

- Uniform-cost expands equally in all “directions”
- A* expands mainly toward the goal, but does hedge its bets to ensure optimality
- Start
- Goal
- Start
- Goal

Converted shape

Start

Converted shape

Goal

Video of Demo Contours (Empty) -- UCS

Video of Demo Contours (Empty) -- Greedy

Video of Demo Contours (Empty) – A*

**Video of Demo Contours (Pacman Small Maze) -
A***

Comparison

- Greedy
- Uniform Cost
- A*



A* Applications



Video of Demo Pacman (Tiny Maze) - UCS / A*

**Video of Demo Empty Water Shallow/Deep -
Guess Algorithm**

Creating Heuristics



Example: 8 Puzzle

- What are the states?
- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?
- Start State
- Goal State
- Actions

Converted shape

Group of shapes

8 Puzzle I

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) =$
- This is a relaxed-problem heuristic
- 8
- Statistics from Andrew Moore

Converted shape

8

8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?
- Total Manhattan distance
- Why is it admissible?
- $h(\text{start}) =$
- $3 + 1 + 2 + \dots = 18$

Converted shape

$$3 + 1 + 2 + \dots = 18$$

Converted shape

Start State Goal State

8 Puzzle III

- How about using the actual cost as a heuristic?
- Would it be admissible?
- Would we save on nodes expanded?
- What's wrong with it?
- With A*: a trade-off between quality of estimate and work per node
- As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself



Semi-Lattice of Heuristics

Trivial Heuristics, Dominance

- Dominance: $h_a \geq h_c$ if
- Heuristics form a semi-lattice:
- Max of admissible heuristics is admissible
- Trivial heuristics
- Bottom of lattice is the zero heuristic (what does this give us?)
- Top of lattice is the exact heuristic

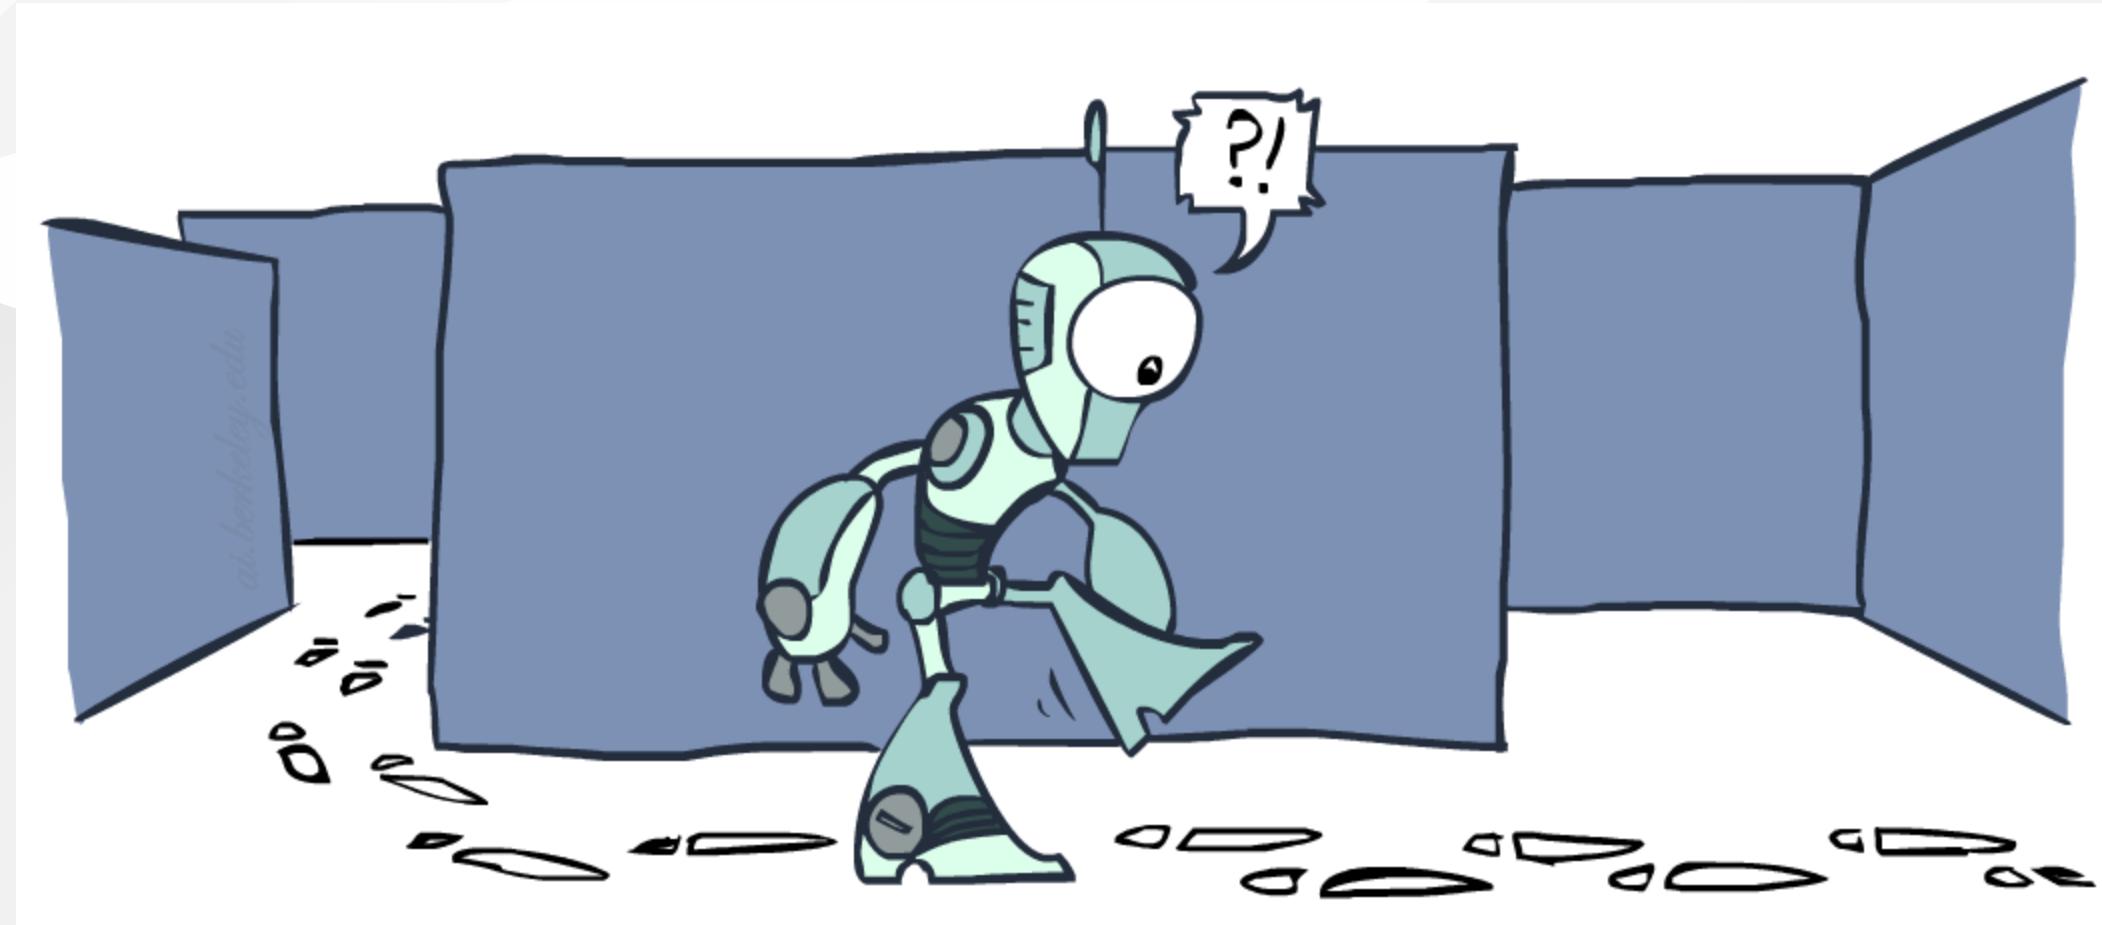
$$h(n) = \max(h_a(n), h_b(n))$$

$$\forall n : h_a(n) \geq h_c(n)$$

Converted shape

Group of shapes

Graph Search



BFS Graph Search Example

- we shouldn't bother expanding the circled nodes: WHY?

Converted shape

S a b d p a c q

Graph Search

- Idea: never expand a state twice
- How to implement:
- Tree search + set of expanded states ("closed set")
- Expand the search tree node-by-node, but...
- Before expanding a node, check to make sure its state has never been expanded before
- If not new, skip it, if new add to closed set
- Important: store the closed set as a set, not a list
- Can graph search wreck completeness? Why/why not?
- How about optimality?

A* Graph Search Gone Wrong?

- S
- A
- B
- C
- G
- 1
- 1
- 1
- 2
- 3
- S (0+2)
- State space graph
- Search tree

Consistency of Heuristics

- Main idea: estimated heuristic costs \leq actual costs
- Admissibility: heuristic cost \leq actual cost to goal
- $h(A) \leq$ actual cost from A to G
- Consistency: heuristic “arc” cost \leq actual cost for each arc
- $h(A) - h(C) \leq \text{cost}(A \text{ to } C)$
- Consequences of consistency:
 - The f value along a path never decreases
 - $h(A) \leq \text{cost}(A \text{ to } C) + h(C)$
 - A* graph search is optimal
- 3
- A
- C
- G

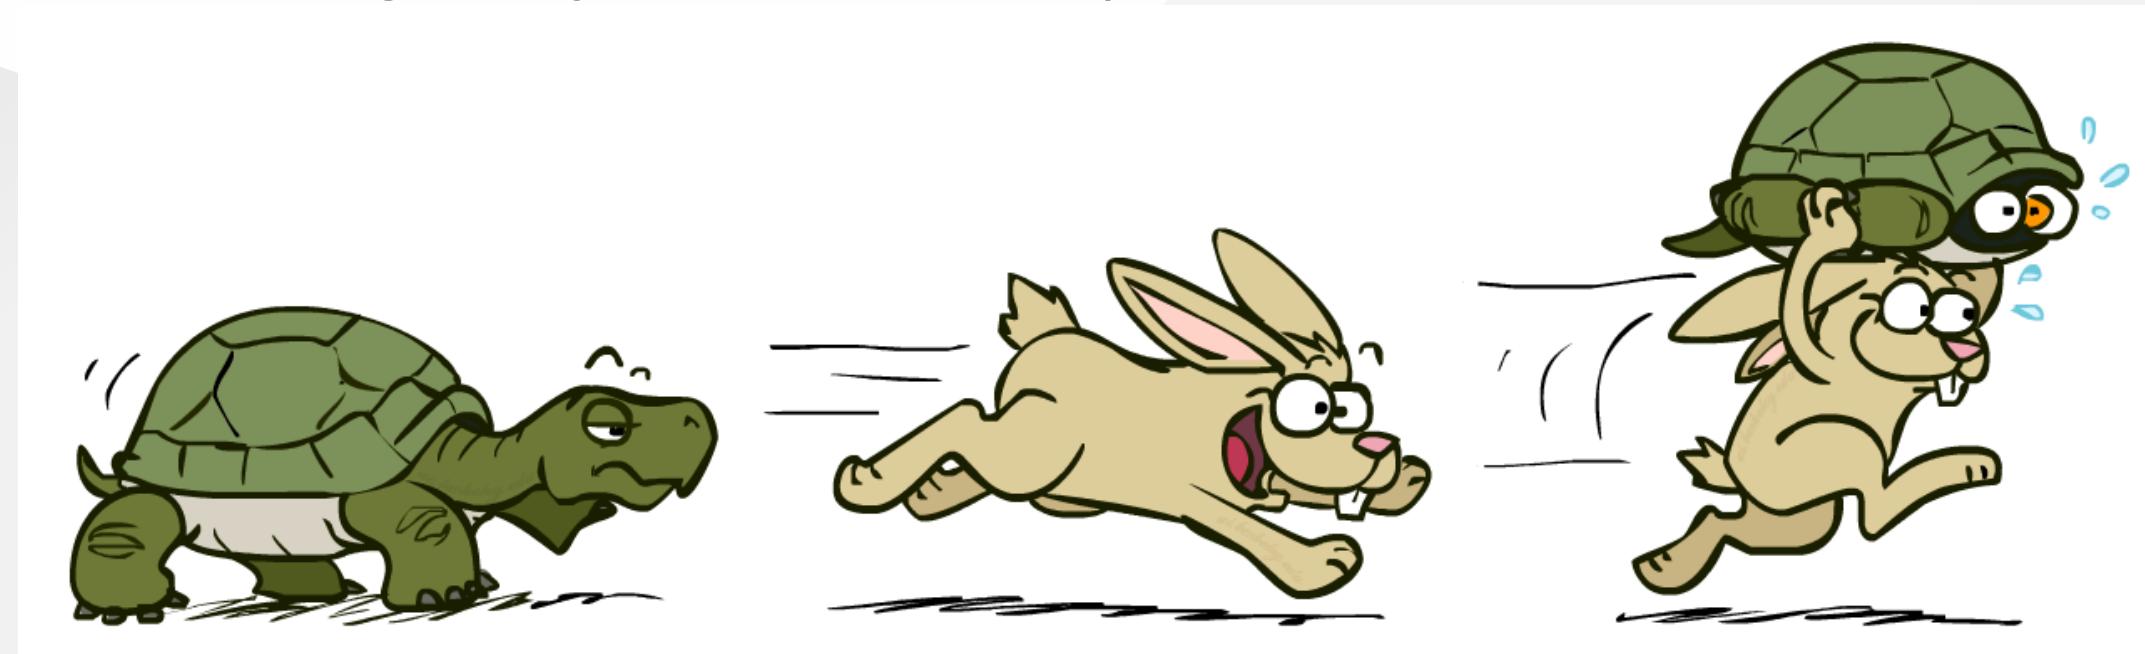
Optimality

- Tree search:
 - A* is optimal if heuristic is admissible
 - UCS is a special case ($h = 0$)
- Graph search:
 - A* optimal if heuristic is consistent
 - UCS optimal ($h = 0$ is consistent)
 - Consistency implies admissibility
- In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems



A*: Summary

- A* uses both backward costs and (estimates of) forward costs
- A* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems



Tree Search Pseudo-Code

```
function TREE-SEARCH(problem, fringe) return a solution, or failure
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    for child-node in EXPAND(STATE[node], problem) do
      fringe  $\leftarrow$  INSERT(child-node, fringe)
    end
  end
```

Graph Search Pseudo-Code

```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed  $\leftarrow$  an empty set
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
```